

Exercise 1

Due on: Thursday, 25.04.2024

Task 1 Tensors

Generate the following matrix without using Python loops:

$$\begin{pmatrix} 4 & 5 & 4 & 4 & 4 & 4 & 5 & 4 & 4 & 4 & 4 & 5 & 4 \\ 5 & 5 & 5 & 5 & 5 & 5 & 5 & 5 & 5 & 5 & 5 & 5 & 5 \\ 4 & 5 & 4 & 4 & 4 & 4 & 5 & 4 & 4 & 4 & 4 & 5 & 4 \\ 4 & 5 & 4 & 6 & 6 & 4 & 5 & 4 & 6 & 6 & 4 & 5 & 4 \\ 4 & 5 & 4 & 6 & 6 & 4 & 5 & 4 & 6 & 6 & 4 & 5 & 4 \\ 4 & 5 & 4 & 4 & 4 & 4 & 5 & 4 & 4 & 4 & 4 & 5 & 4 \\ 5 & 5 & 5 & 5 & 5 & 5 & 5 & 5 & 5 & 5 & 5 & 5 & 5 \\ 4 & 5 & 4 & 4 & 4 & 4 & 5 & 4 & 4 & 4 & 4 & 5 & 4 \\ 4 & 5 & 4 & 6 & 6 & 4 & 5 & 4 & 6 & 6 & 4 & 5 & 4 \\ 4 & 5 & 4 & 6 & 6 & 4 & 5 & 4 & 6 & 6 & 4 & 5 & 4 \\ 4 & 5 & 4 & 4 & 4 & 4 & 5 & 4 & 4 & 4 & 4 & 5 & 4 \\ 5 & 5 & 5 & 5 & 5 & 5 & 5 & 5 & 5 & 5 & 5 & 5 & 5 \\ 4 & 5 & 4 & 4 & 4 & 4 & 5 & 4 & 4 & 4 & 4 & 5 & 4 \end{pmatrix}$$

Hint: use `torch.full` and the slicing operator

Task 2 Softmax Function

Let $\mathbf{x} \in \mathbb{R}^d$ be a vector. The softmax function $\sigma : \mathbb{R}^d \rightarrow [0, 1]^d$ is given by

$$\sigma(\mathbf{x})_j = \frac{\exp(x_j)}{\sum_{k=1}^d \exp(x_k)} \quad \text{for } j = 1, \dots, d.$$

Implement the sigmoid and the softmax functions in PyTorch and compare their outputs for a random vector. What can you say about the highest values, respectively?

Task 3 Empirical Verification of Perceptron Convergence

Consider a perceptron $f(x) = \text{sign}(w \cdot x + b)$ and the following training algorithm:

- (1) Initialize w^0 , i.e., via $w^0 = \mathbf{0}$ (or randomly).
- (2) While there is some $n_k \in \{1, \dots, n\}$ such that $y_{n_k}(w^k \cdot x_{n_k}) \leq 0$,
update $w^{k+1} = w^k + y_{n_k} x_{n_k}$.

Note that this assumes that b is added into w and that x is augmented with a 1.

Assume for every n that:

- (1) $\|x_n\|_2 < R$, with some $R > 0$.
- (2) There exists a margin $\gamma > 0$ and some non-zero w^* (with $\|w^*\|_2 = 1$) such that $y_n(w^* \cdot x_n) \geq \frac{\gamma}{2}$.

Confirm empirically that training this model will converge in $k \leq \frac{4R^2}{\gamma^2}$ updates by following these steps:

- (i) Create a linearly separable toy data set.
- (ii) Compute the R and γ values for that data set.
- (iii) Implement the perceptron algorithm.
- (iv) Initialize the perceptron at least 250 times with random weights, track the number of updates, and confirm that those numbers are smaller than the bound.

Hint: Take a look into `sklearn.datasets` for creating a toy data set. You can compute the margin using an SVM, i.e., `sklearn.svm.LinearSVC`. Experiment with tighter and wider margins. There is a template `code1.py`.

Task 4 The Perceptron on MNIST

- (i) Load the MNIST data set, i.e., via `torchvision.datasets.MNIST`
- (ii) Visualize some of the numbers.
- (iii) Restrict the multi-class classification problem to a binary classification of digit 0 vs digit 1.
- (iv) Run the perceptron algorithm on this data set.
- (v) Compute the train and test accuracies.
- (vi) Show the misclassified input instances.

Hint: Take a look into `code1.py`.