

## Exercise 2

Due on: Thursday, 02.05.2024

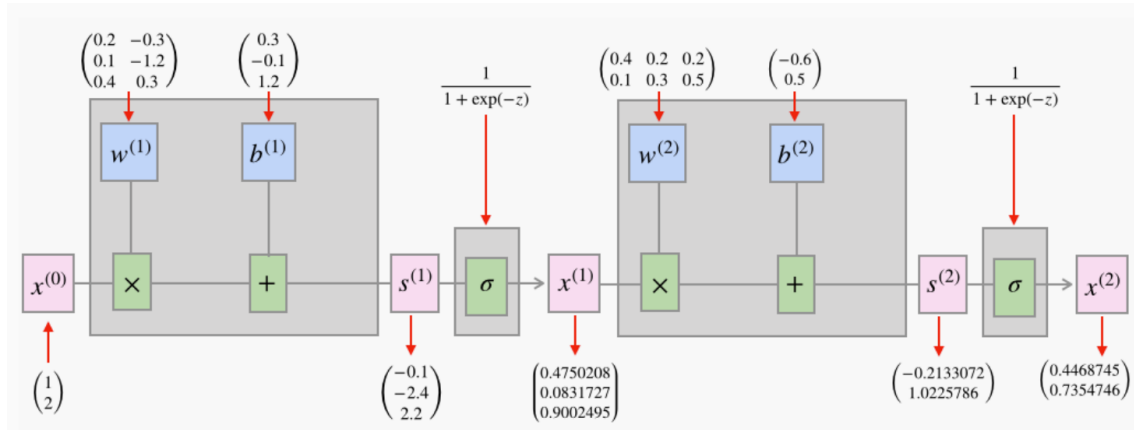


Figure 1: A simple neural network.

### Task 5 A Simple Neural Network

Let  $x^{(0)} = (1, 2)^\top \in \mathbb{R}^2$  be a data point and  $y = (0, 1)^\top \in \mathbb{R}^2$  be the target. Consider the neural network with sigmoid activation functions depicted in Figure 1.

- Compute the forward pass of  $x^{(0)}$  by hand. What is the value of  $x^{(2)}$ ?
- Compute the gradients with respect to all parameters in the network using backpropagation. For simplicity, you can use the squared loss function.

### Solution

- Figure 1 now contains the (intermediate) result(s).
- Let  $E(\hat{y}, y) = \frac{1}{2}(\hat{y} - y)^2$  be the squared loss function. Its derivative is given by  $\frac{\partial E}{\partial \hat{y}} = \hat{y} - y$ . Here,  $\hat{y}$  is the prediction, hence  $x^{(2)}$  in our case. Remember, that for a sigmoid activation function, the derivative is  $\sigma'(\cdot) = \sigma(\cdot)(1 - \sigma(\cdot))$ .

$$\begin{aligned}
\frac{\partial E}{\partial w^{(2)}} &= \underbrace{\frac{\partial E}{\partial x^{(2)}} \frac{\partial x^{(2)}}{\partial s^{(2)}}}_{\text{tmp}} \frac{\partial s^{(2)}}{\partial w^{(2)}} \\
&= \left( x^{(2)} - y \right) \sigma \left( s^{(2)} \right) \left( 1 - \sigma \left( s^{(2)} \right) \right) x^{(1)\top} \\
&= \begin{pmatrix} 0.45 \\ -0.26 \end{pmatrix} \begin{pmatrix} 0.25 \\ 0.19 \end{pmatrix} \begin{pmatrix} 0.48 & 0.08 & 0.90 \end{pmatrix} \\
&= \begin{pmatrix} 0.05 & 0.01 & 0.10 \\ -0.02 & -0.00 & -0.05 \end{pmatrix} \\
\frac{\partial E}{\partial b^{(2)}} &= \underbrace{\frac{\partial E}{\partial x^{(2)}} \frac{\partial x^{(2)}}{\partial s^{(2)}}}_{\text{tmp}} \frac{\partial s^{(2)}}{\partial b^{(2)}} \\
&= \left( x^{(2)} - y \right) \sigma \left( s^{(2)} \right) \left( 1 - \sigma \left( s^{(2)} \right) \right) 1 \\
&= \begin{pmatrix} 0.45 \\ -0.26 \end{pmatrix} \begin{pmatrix} 0.25 \\ 0.19 \end{pmatrix} \\
&= \begin{pmatrix} 0.11 \\ -0.05 \end{pmatrix} \\
\frac{\partial E}{\partial w^{(1)}} &= \underbrace{\frac{\partial E}{\partial x^{(2)}} \frac{\partial x^{(2)}}{\partial s^{(2)}}}_{\text{previous tmp}} \frac{\partial s^{(2)}}{\partial x^{(1)}} \frac{\partial x^{(1)}}{\partial s^{(1)}} \frac{\partial s^{(1)}}{\partial w^{(1)}} \\
&= \underbrace{\text{tmp} \frac{\partial s^{(2)}}{\partial x^{(1)}} \frac{\partial x^{(1)}}{\partial s^{(1)}}}_{\text{new tmp}} \frac{\partial s^{(1)}}{\partial w^{(1)}} \\
&= \text{tmp} w^{(2)\top} \sigma \left( s^{(1)} \right) \left( 1 - \sigma \left( s^{(1)} \right) \right) x^{(0)\top} \\
&= w^{(2)\top} \text{tmp} \sigma \left( s^{(1)} \right) \left( 1 - \sigma \left( s^{(1)} \right) \right) x^{(0)\top} \\
&= \begin{pmatrix} 0.4 & 0.1 \\ 0.2 & 0.3 \\ 0.2 & 0.5 \end{pmatrix} \begin{pmatrix} 0.11 \\ -0.05 \end{pmatrix} \begin{pmatrix} 0.25 \\ 0.08 \\ 0.09 \end{pmatrix} \begin{pmatrix} 1.00 & 2.00 \end{pmatrix} \\
&= \begin{pmatrix} 0.010 & 0.019 \\ 0.001 & 0.001 \\ 0.000 & -0.001 \end{pmatrix} \\
\frac{\partial E}{\partial b^{(1)}} &= \underbrace{\frac{\partial E}{\partial x^{(2)}} \frac{\partial x^{(2)}}{\partial s^{(2)}}}_{\text{previous tmp}} \frac{\partial s^{(2)}}{\partial x^{(1)}} \frac{\partial x^{(1)}}{\partial s^{(1)}} \frac{\partial s^{(1)}}{\partial b^{(1)}} \\
&= \frac{\partial E}{\partial x^{(2)}} \frac{\partial x^{(2)}}{\partial s^{(2)}} \frac{\partial s^{(2)}}{\partial x^{(1)}} \frac{\partial x^{(1)}}{\partial s^{(1)}} 1 \\
&= \begin{pmatrix} 0.010 \\ 0.001 \\ 0.000 \end{pmatrix}
\end{aligned}$$

## Task 6 A Neural Network Implementation

We now want to implement a neural network as a class in Python. Check the template `code2.py`.

### Part 1

- (i) Implement the ReLU and sigmoid activation functions and their derivatives.
- (ii) The class `NeuralNetwork` already has two member functions `__init__` and `forward`. Fix the TODOs and
  - (a) initialize the weights via  $w_{ij} \sim \mathcal{N}(0, \sigma^2)$ , where  $\sigma = \sqrt{\frac{2}{d_{\text{in}} + d_{\text{out}}}}$ ,  $d_{\text{in}}$  denotes the number of input neurons, and  $d_{\text{out}}$  denotes the number of output neurons of that layer
  - (b) initialize the biases as zeros
  - (c) implement the forward pass

Make sure to also save the intermediate results (e.g.  $x^{(\ell)}$ ,  $s^{(\ell)}$ , see Figure 1) when computing the forward pass. We will need those later on.

*Hint:* You can test your implementation by overwriting the weights and biases with the values from Task 5 (i) and compare all intermediate results.

### Part 2

- (i) Implement backpropagation as a member function `backprop(self, x, y)`. It should compute the gradients of all parameters.
- (ii) Implement a training function `train(self, x, y, eta=.01, iterations=100)` that implements (batch) gradient descent. You can use the squared loss if not stated otherwise.
- (iii) Test your training procedure by fitting a binary classification problem on toy data and visualize the prediction appropriately. Make sure to use a sigmoid activation and a single neuron in the output layer.

*Hint:* You can test your implementation by comparing the results with the values from Task 5 (ii).

### Solution

The code is provided as `solution2.py`.

## Task 7 Approximating a Function Using ReLUs

In this task, we aim to approximate a function  $f : [a, b] \rightarrow \mathbb{R}$  arbitrarily well using a neural network with a single hidden layer consisting of  $n$  neurons based on ReLU activations.

Consider a grid of  $n$  equidistant points  $x_1, x_2, x_3, \dots, x_n$ , with  $a = x_1$  and  $b = x_n$ . Furthermore, the grid has a constant spacing of  $h = x_{i+1} - x_i$  for  $i = 1, 2, \dots, n-1$ . We want to approximate  $f$  using a linear combination of ReLU functions given by

$$\text{NN}(x) = \sum_{i=1}^n \lambda_i \cdot \text{ReLU}(x - b_i),$$

where  $\lambda_i \in \mathbb{R}$  are weights and the biases are given by  $b_1 = a - h$ ,  $b_2 = a$ ,  $b_3 = a + h$ ,  $\dots$ ,  $b_n = b - h$ .

Prior to  $b_1$ , all ReLU functions  $\max(0, x - b_i)$  are zero. Between  $b_1 = a - h$  and  $b_2 = a = x_1$ , only the first ReLU function  $\max(0, x - b_1)$  can be non-zero. The function  $\lambda_1 \cdot \max(0, x - b_1)$  is supposed to approximate  $f$  at  $x_1$ . If we subtract that function from  $f$ , we obtain a new function which the next ReLU function can approximate, and iterate to the end. Hence, we can approximate the function  $f$  on the interval  $[a, b]$  using a sum of  $n$  ReLU functions.

To learn the parameters  $\lambda_i$  we can construct a linear system of equations

$$f(x_i) = \sum_{j=1}^n \lambda_j \cdot \text{ReLU}(x_i - b_j) \quad \text{for } i = 1, \dots, n.$$

By using  $x_i = a + (i-1)h$  and  $b_j = a + (j-2)h$ , we obtain  $x_i - b_j = (i-j+1)h$ . Considering  $i$  and  $j$  where the ReLU is positive,

$$f(x_i) = \sum_{j=1}^i \lambda_j (i-j+1)h \quad \text{for } i = 1, \dots, n. \quad (1)$$

Write down the linear system determined by Equation (1). Then, implement this approximation method, solve for  $\lambda_1, \lambda_2, \dots, \lambda_n$ , and experiment with various values of  $n$ . Check `code2.py`.

## Solution

Equation (1) yields

$$\begin{aligned} f(x_1) &= 1h\lambda_1 &= h(1\lambda_1) \\ f(x_2) &= 2h\lambda_1 + 1h\lambda_2 &= h(2\lambda_1 + 1\lambda_2) \\ f(x_3) &= 3h\lambda_1 + 2h\lambda_2 + 1h\lambda_3 &= h(3\lambda_1 + 2\lambda_2 + 1\lambda_3) \\ f(x_4) &= 4h\lambda_1 + 3h\lambda_2 + 2h\lambda_3 + 1h\lambda_4 &= h(4\lambda_1 + 3\lambda_2 + 2\lambda_3 + 1\lambda_4) \end{aligned}$$

for the first values of  $i$  which can be written in the following linear system of equations:

$$h \begin{pmatrix} 1 & & & \\ 2 & 1 & & \\ \vdots & \ddots & \ddots & \\ n & \dots & 2 & 1 \end{pmatrix} \begin{pmatrix} \lambda_1 \\ \lambda_2 \\ \vdots \\ \lambda_n \end{pmatrix} = \begin{pmatrix} f_1 \\ f_2 \\ \vdots \\ f_n \end{pmatrix}$$

where  $f_i = f(x_i)$  for  $i = 1, \dots, n$ . The parameters  $\lambda_i$  can be obtained by solving the system. Note that this approach can be generalized to non-equidistant data points. The approximation improves by increasing the number of functions  $n$ .

The code is provided as `solution2.py`.